Information Arrival in Financial Markets

Michael G. Zdinak

Washington University in St. Louis

DOWNLOAD LATEST VERSION

November 14, 2019

Abstract

This article introduces a new high-frequency analysis of six years of data for options written on the S&P 500 and traded on the Chicago Board of Exchange. I quantify in real-time the information contained in the probability measure implied by option prices, using concepts developed in information theory. Here information is analogous to a reduction in uncertainty surrounding the future price of the underlying security. A simple nonparametric estimator allows us to measure the amount of information gained as an option approaches maturity. I then test for jumps in the expectation of said future price. I find the intraday flow of information in a large and important market is not continuous, and often increases in discrete intervals. This fact is used to identify events in which a large amount of information is revealed to investors.

Introduction

How does information arrive in financial markets? In this paper, I confront the basic characterization of the process by which investors learn about the future value of an asset. The topic is of importance to much of financial economics, yet continues to be one of the least explored. Indeed, this paper is the first to quantify in real time how information drives price discovery in option markets. The results document two new facts. First, the majority of information concerning the future price of an asset accrues only in the final month of an option's life. Second, information often arrives in discrete intervals resulting in jumps in investor expectations. These innovations represent valuable insights to the field of financial economics. In both cases, existing theoretical models can be shown to not accurately describe either process.

The paper joins a growing literature of high-frequency analysis of investor expectations, of which Birru and Figlewski (2010) offer another example. Both the literature and this paper estimate the density of future returns as implied by observed option prices. Following Cox and Ross (1976) and Cox, Ross, and Rubinstein (1979), this estimation relies on a representative investor's ability to arbitrage an option and its underlying asset. In that way, all risk except the underlying uncertainty surrounding the asset's future price may be hedged away. The resulting implied distribution of future returns is known as the "risk-neutral density". Figlewski (2017) offers a recent review of the key ideas in the literature. Following Harrison and Kreps (1979), if this density is known, then options may be priced as if all investors are "risk-neutral". Option prices therefore reflect investor beliefs over the probability the world will achieve some future state; and it is this information that is of vital interest to investors, researchers, and policymakers. This fact combined with the growth in the derivatives market has inspired renewed interest in understanding how investors respond to new information.

In this paper, I provide a high-frequency analysis of the price discovery process in option markets. Using six years of data for options written on the S&P 500 and traded on the Chicago Board of Exchange, I characterize the evolution of the density function for options with the same maturity date. This paper is the first to estimate the intraday dynamics of the risk-neutral density over an option's life-cycle, and offers the following three methodological contributions to the literature. First, I show how a simple nonparametric estimator can be used to approximate the implied density of future returns at high-frequencies. Second, I show how concepts developed in information theory can be used to quantify the amount of information contained in the estimated density. Information is then a reduction in uncertainty surrounding the future price of the underlying security. Third, I show that this novel approach permits a simple testing procedure for the presence of jumps in the arrival of information and the evolution of the risk-neutral density. The results of this testing procedure represent the paper's main contribution to the literature.

I find that information often arrives in discrete intervals. Even at high-frequencies the risk-neutral density can be shown to "jump". The testing procedure reveals both the frequency and magnitude of the jumps in investor expectations. I identify at least one jump for a majority of days, and find days without jumps contribute little to the total information gained over an option's life. I then document two empirical facts new to the literature. First, the majority of information accrues in the final month of an option's life-cycle. In fact, I show investors learn little about the future price of an asset for most of an option's life. Second, jumps contribute the majority of information gained during the first two months. Only in the final month does information arrives often enough to contribute more to the total gained. The size and frequency of jumps does not decline.

The paper builds on earlier work in many ways, but several features distinguish the findings from previous results. These include; 1. a focus on the evolution of the risk-neutral density over an option's life cycle, 2. a fully nonparametric estimation technique, 3. a measure of information as a reduction in uncertainty, 4. a simple framework to test for jumps, and 5. the frequency and length of the sample of options data. These features are discussed briefly in turn.

This paper targets the evolution of the risk-neutral density. In a departure from much of the earlier literature (Ait-Sahalia and Lo (1998), among others), the density is not estimated for a fixed maturity. Instead, for options with the same expiration date, the focus is on how the density becomes more concentrated as the settlement date approaches. The perspective permits a simple testing procedure for the presence of jumps in the evolution of the density, and does not require the estimated density to be interpolated over time.

The paper employs an estimation procedure that is fully nonparametric. Inspired by the original procedure proposed by Breeden and Litzenberger (1978), the estimator places no restrictions on the shape of the density function or the dynamics of underlying asset's price. In contrast to other semiparametric procedures, such as those following Shimko (1993), the estimation does not require the interpolation or the extrapolation of observed option prices, or the need for significant tradeoffs in measures of the goodness-of-fit and smoothness.

This paper introduces concepts from information theory to quantify the uncertainty investors face about the future value of the underlying asset. The basic insight is that information can be measured as the reduction in uncertainty. The basic quantity of information theory is entropy. The concept is not new to economics, both Sims (2003) and Frankel and Kamenica (2019) use entropy to model rational inattention and quantify the information in a decision problem. Stutzer (2000) and Buchen and Kelly (1996) go so far as to use the concept of maximum entropy to infer the risk-neutral density from observed prices. However, this paper represents the first application of entropy to the problem of quantifying the information in the density function over time.

The paper employs a framework and hypothesis test for jumps in the risk-neutral density. The evolution of the entropy of the risk-neutral density reflects the arrival of information. The procedure is equivalent to testing for jumps in a nonstationary timeseries, following Zivot and Andrews (1992). The test derives from the literature on testing for unit roots in economic timeseries, beginning with Dickey and Fuller (1979), and generalized by Said and Dickey (1984).

The choice of procedure is threefold. First the test is simple and transparent. It is fast and easy to implement, and the results are easily interpreted. Second, the procedure reveals the frequency, magnitude, and timing of each jump, three items that are of immediate interest. Third, the procedure permits a statistical test for each identified jump. The analysis therefore differs in important respects from even high-frequency event-studies, such as Goldberg and Grisse (2013) and Andersen, Bollerslev, Diebold and Vega (2003), among others, who examine the response of interest rates and exchange rates to economic news announcements.

Finally, the paper uses a new dataset of intraday quotes for all options written on the S&P 500 and traded on the Chicago Board of Exchange. The data is novel in two respects, namely the high frequency and long calendar span of the sample of SPX options analyzed. The analysis covers six years, or nearly 1,500 trading days, beginning in January 2009 and ending in December 2014. In comparison to other high-frequency studies, Jiang and Tian (2005) and Birru and Figlewski (2010), which focus on forecasting realized variance and the 1-minute changes in the quantiles of risk-neutral density during the fall of 2008, the data presented here represents the most complete picture of the intraday evolution investor expectations found in the literature to date.

The paper proceed as follows. In section 2, I describe the high-frequency options data. Section 3 characterizes the nonparametric estimator used to approximate the risk-neutral density.

In section 4 I propose the use of information theory to quantify the information in the estimated density. Section 5 describes the framework to test for jumps in the arrival of information. Section 6 discusses the results of the hypothesis tests. In section 7 I present select case studies around events where there was a large jump in information. Finally, section 8 concludes.

Selected Bibliography

Aït-sahalia, Y., & Lo, A. W. (1998). Nonparametric Estimation of State-Price. *The Journal of Finance*, *53*(2), 499–547.

Andersen, T. G., Bollerslev, T., Diebold, F. X., & Vega, C. (2003). Micro effects of macro announcements: Real-time price discovery in foreign exchange. *American Economic Review*, *93*(1), 38–62. https://doi.org/10.1257/000282803321455151

Bahra, B. (2007). Implied risk-neutral probability density functions from option prices: A central bank perspective. *Forecasting Volatility in the Financial Markets*, 201–226. https://doi.org/10.1016/B978-075066942-9.50011-X

Bennett, D. L., & Gourley, S. A. (1953). sharing partner libraries or were purchased by Olin Library Contact Interlibrary Loan Research Question ? *Applied Mathematics and Computation*, *39*(4), 198–201. https://doi.org/10.1300/J108v05n03

Birru, J., & Figlewski, S. (2012). Anatomy of a meltdown: The risk neutral density for the S&P 500 in the fall of 2008. *Journal of Financial Markets*, *15*(2), 151–180. https://doi.org/10.1016/j.finmar.2011.09.001

Bollerslev, T., Russell, J. R., Oxford, M. W., Oxford, U. K., & Figlewski, S. (2008). Estimating the Implied Risk Neutral Density for the U . S . Market Portfolio by The Implied Risk Neutral Density for the U . S . Market Portfolio Abstract. *Portfolio The Magazine Of The Fine Arts*.

Breeden, D. T., & Litzenberger, R. H. (1978). Prices of State-Contingent Claims Implicit in Option Prices. *The Journal of Business*, *51*(4), 621. https://doi.org/10.1086/296025

Buchen, P. W., & Kelly, M. (2016). The Maximum Entropy Distribution of an Asset Inferred from Option Prices Author (s): Peter W . Buchen and Michael Kelly Source : The Journal of Financial and Quantitative Analysis , Vol . 31 , No . 1 (Mar ., 1996), pp . Published by : Cambridge Univer, *31*(1), 143–159.

Datta, D. D., Londono, J. M., & Ross, L. J. (2017). Generating options-implied probability densities to understand oil market events. *Energy Economics*, *64*(1122), 440–457. https://doi.org/10.1016/j.eneco.2016.01.006

Figlewski, S. (2018). Risk Neutral Densities : A Review Risk Neutral Densities : A Review.

Frankel, A., & Kamenica, E. (2019). Quantifying Information and Uncertainty. *American Economic Review*, *109*(9), 3650–3680. https://doi.org/10.1257/aer.20181897

Goldberg, L. S., & Grisse, C. (2013). Time Variation in Asset Price Responses to Macro Announcements. *SSRN Electronic Journal*, (626). https://doi.org/10.2139/ssrn.2321203

Jiang, G. J., & Tian, Y. S. (2005). The model-free implied volatility and its information content. *Review of Financial Studies*, *18*(4), 1305–1342. https://doi.org/10.1093/rfs/hhi027

Shimko, D. (1993). Bounds of probability. *Risk Publications*, *6*(4), 33–37. Retrieved from https://www.researchgate.net/publication/306151578_Bounds_of_probability

Sims, C. A. (2003). Implications of rational inattention. *Journal of Monetary Economics*, *50*(3), 665–690. https://doi.org/10.1016/S0304-3932(03)00029-1

Stutzer, M. J. (2000). Simple entropic derivation of a generalized black-scholes option pricing model. *Entropy*, *2*(2), 70–77. https://doi.org/10.3390/e2020070

Zivot, E., & Andrews, D. W. K. (1992). Further evidence on the great crash, the oil-price shock, and the unit-root hypothesis. *Journal of Business and Economic Statistics*, *10*(3), 251–270. https://doi.org/10.1080/07350015.1992.10509904